

# tf.estimator Quickstart

TensorFlow’s high-level machine learning API (tf.estimator) makes it easy to configure, train, and evaluate a variety of machine learning models. In this tutorial, you’ll use tf.estimator to construct a [neural network](https://en.wikipedia.org/wiki/Artificial_neural_network) (https://en.wikipedia.org/wiki/Artificial\_neural\_network) classifier and train it on the [Iris data set](https://en.wikipedia.org/wiki/Iris_flower_data_set) (https://en.wikipedia.org/wiki/Iris\_flower\_data\_set) to predict flower species based on sepal/petal geometry. You'll write code to perform the following five steps:

- 1. Load CSVs containing Iris training/test data into a TensorFlow **Dataset**
- 2. Construct a [neural network classifier](https://www.tensorflow.org/api_docs/python/tf/estimator/DNNClassifier) (https://www.tensorflow.org/api\_docs/python/tf/estimator/DNNClassifier?hl=zh-cn)
- 3. Train the model using the training data
- 4. Evaluate the accuracy of the model
- 5. Classify new samples

NOTE: Remember to [install TensorFlow on your machine](https://www.tensorflow.org/install/index?hl=zh-cn) (https://www.tensorflow.org/install/index?hl=zh-cn) before getting started with this tutorial.

## Complete Neural Network Source Code

Here is the full code for the neural network classifier:

```
from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

import os
from six.moves.urllib.request import urlopen

import numpy as np
import tensorflow as tf

# Data sets
IRIS_TRAINING = "iris_training.csv"
IRIS_TRAINING_URL = "http://download.tensorflow.org/data/iris_training.csv"

IRIS_TEST = "iris_test.csv"
IRIS_TEST_URL = "http://download.tensorflow.org/data/iris_test.csv"

def main():
    # If the training and test sets aren't stored locally, download them.
    if not os.path.exists(IRIS_TRAINING):
        raw = urlopen(IRIS_TRAINING_URL).read()
        with open(IRIS_TRAINING, "wb") as f:
            f.write(raw)

    if not os.path.exists(IRIS_TEST):
        raw = urlopen(IRIS_TEST_URL).read()
        with open(IRIS_TEST, "wb") as f:
            f.write(raw)

    # Load datasets.
    training_set = tf.contrib.learn.datasets.base.load_csv_with_header(
        filename=IRIS_TRAINING,
        target_dtype=np.int,
        features_dtype=np.float32)
    test_set = tf.contrib.learn.datasets.base.load_csv_with_header(
        filename=IRIS_TEST,
        target_dtype=np.int,
        features_dtype=np.float32)

    # Specify that all features have real-value data
    feature_columns = [tf.feature_column.numeric_column("x", shape=[4])]

    # Build 3 layer DNN with 10, 20, 10 units respectively.
    classifier = tf.estimator.DNNClassifier(feature_columns=feature_columns,
```

```

                                hidden_units=[10, 20, 10],
                                n_classes=3,
                                model_dir="/tmp/iris_model")

# Define the training inputs
train_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={"x": np.array(training_set.data)},
    y=np.array(training_set.target),
    num_epochs=None,
    shuffle=True)

# Train model.
classifier.train(input_fn=train_input_fn, steps=2000)

# Define the test inputs
test_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={"x": np.array(test_set.data)},
    y=np.array(test_set.target),
    num_epochs=1,
    shuffle=False)

# Evaluate accuracy.
accuracy_score = classifier.evaluate(input_fn=test_input_fn) ["accuracy"]

print("\nTest Accuracy: {0:f}\n".format(accuracy_score))

# Classify two new flower samples.
new_samples = np.array(
    [[6.4, 3.2, 4.5, 1.5],
     [5.8, 3.1, 5.0, 1.7]], dtype=np.float32)
predict_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={"x": new_samples},
    num_epochs=1,
    shuffle=False)

predictions = list(classifier.predict(input_fn=predict_input_fn))
predicted_classes = [p["classes"] for p in predictions]

print(
    "New Samples, Class Predictions:    {}\n"
    .format(predicted_classes))

if __name__ == "__main__":
    main()
```

The following sections walk through the code in detail.

## Load the Iris CSV data to TensorFlow

The [Iris data set](https://en.wikipedia.org/wiki/Iris_flower_data_set) ([https://en.wikipedia.org/wiki/Iris\\_flower\\_data\\_set](https://en.wikipedia.org/wiki/Iris_flower_data_set)) contains 150 rows of data, comprising 50 samples from each of three related Iris species: *Iris setosa*, *Iris virginica*, and *Iris versicolor*.



From left to right, ***Iris setosa*** (<https://commons.wikimedia.org/w/index.php?curid=170298>) (by **Radomil** (<https://commons.wikimedia.org/wiki/User:Radomil>), **CC BY-SA 3.0**), ***Iris versicolor*** (<https://commons.wikimedia.org/w/index.php?curid=248095>) (by **Dlanglois** (<https://commons.wikimedia.org/wiki/User:Dlanglois>), **CC BY-SA 3.0**), and ***Iris virginica*** (<https://www.flickr.com/photos/33397993@N05/3352169862>) (by **Frank Mayfield** (<https://www.flickr.com/photos/33397993@N05>), **CC BY-SA 2.0**).

Each row contains the following data for each flower sample: sepal (<https://en.wikipedia.org/wiki/Sepal>) length, sepal width, petal (<https://en.wikipedia.org/wiki/Petal>) length, petal width, and flower species. Flower species are represented as integers, with 0 denoting *Iris setosa*, 1 denoting *Iris versicolor*, and 2 denoting *Iris virginica*.

Sepal Length	Sepal Width	Petal Length	Petal Width	Species
5.1	3.5	1.4	0.2	0
4.9	3.0	1.4	0.2	0
4.7	3.2	1.3	0.2	0
...	...	...	...	...
7.0	3.2	4.7	1.4	1
6.4	3.2	4.5	1.5	1
6.9	3.1	4.9	1.5	1
...	...	...	...	...
6.5	3.0	5.2	2.0	2
6.2	3.4	5.4	2.3	2
5.9	3.0	5.1	1.8	2

For this tutorial, the Iris data has been randomized and split into two separate CSVs:

- A training set of 120 samples ([iris\\_training.csv](http://download.tensorflow.org/data/iris_training.csv?hl=zh-cn) ([http://download.tensorflow.org/data/iris\\_training.csv?hl=zh-cn](http://download.tensorflow.org/data/iris_training.csv?hl=zh-cn)))
- A test set of 30 samples ([iris\\_test.csv](http://download.tensorflow.org/data/iris_test.csv?hl=zh-cn) ([http://download.tensorflow.org/data/iris\\_test.csv?hl=zh-cn](http://download.tensorflow.org/data/iris_test.csv?hl=zh-cn))).

To get started, first import all the necessary modules, and define where to download and store the dataset:

```
from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

import os
from six.moves.urllib.request import urlopen

import tensorflow as tf
import numpy as np

IRIS_TRAINING = "iris_training.csv"
IRIS_TRAINING_URL = "http://download.tensorflow.org/data/iris_training.csv"

IRIS_TEST = "iris_test.csv"
IRIS_TEST_URL = "http://download.tensorflow.org/data/iris_test.csv"
```

Then, if the training and test sets aren't already stored locally, download them.

```
if not os.path.exists(IRIS_TRAINING):
    raw = urlopen(IRIS_TRAINING_URL).read()
    with open(IRIS_TRAINING, 'wb') as f:
        f.write(raw)

if not os.path.exists(IRIS_TEST):
    raw = urlopen(IRIS_TEST_URL).read()
    with open(IRIS_TEST, 'wb') as f:
        f.write(raw)
```

Next, load the training and test sets into **Datasets** using the `load_csv_with_header()` (<https://www.github.com/tensorflow/tensorflow/blob/r1.4/tensorflow/contrib/learn/python/learn/datasets/base.py>) method in `learn.datasets.base`. The `load_csv_with_header()` method takes three required arguments:

- **filename**, which takes the filepath to the CSV file
- **target\_dtype**, which takes the numpy datatype (<http://docs.scipy.org/doc/numpy/user/basics.types.html>) of the dataset's target value.

- **features\_dtype**, which takes the numpy datatype (<http://docs.scipy.org/doc/numpy/user/basics.types.html>) of the dataset's feature values.

Here, the target (the value you're training the model to predict) is flower species, which is an integer from 0–2, so the appropriate `numpy` datatype is `np.int`:

```
# Load datasets.
training_set = tf.contrib.learn.datasets.base.load_csv_with_header(
    filename=IRIS_TRAINING,
    target_dtype=np.int,
    features_dtype=np.float32)
test_set = tf.contrib.learn.datasets.base.load_csv_with_header(
    filename=IRIS_TEST,
    target_dtype=np.int,
    features_dtype=np.float32)
```

Datasets in `tf.contrib.learn` are named tuples (<https://docs.python.org/2/library/collections.html#collections.namedtuple>); you can access feature data and target values via the `data` and `target` fields. Here, `training_set.data` and `training_set.target` contain the feature data and target values for the training set, respectively, and `test_set.data` and `test_set.target` contain feature data and target values for the test set.

Later on, in "Fit the DNNClassifier to the Iris Training Data," (`#fit_dnnclassifier`) you'll use `training_set.data` and `training_set.target` to train your model, and in "Evaluate Model Accuracy," (`#evaluate_accuracy`) you'll use `test_set.data` and `test_set.target`. But first, you'll construct your model in the next section.

## Construct a Deep Neural Network Classifier

`tf.estimator` offers a variety of predefined models, called **Estimators**, which you can use "out of the box" to run training and evaluation operations on your data. Here, you'll configure a Deep Neural Network Classifier model to fit the Iris data. Using `tf.estimator`, you can instantiate your `tf.estimator.DNNClassifier` ([https://www.tensorflow.org/api\\_docs/python/tf/estimator/DNNClassifier?hl=zh-cn](https://www.tensorflow.org/api_docs/python/tf/estimator/DNNClassifier?hl=zh-cn)) with just a couple lines of code:

```
# Specify that all features have real-value data
feature_columns = [tf.feature_column.numeric_column("x", shape=[4])]

# Build 3 layer DNN with 10, 20, 10 units respectively.
classifier = tf.estimator.DNNClassifier(feature_columns=feature_columns,
                                       hidden_units=[10, 20, 10],
                                       n_classes=3,
                                       model_dir="/tmp/iris_model")
```

The code above first defines the model's feature columns, which specify the data type for the features in the data set. All the feature data is continuous, so `tf.feature_column.numeric_column` is the appropriate function to use to construct the feature columns. There are four features in the data set (sepal width, sepal height, petal width, and petal height), so accordingly `shape` must be set to `[4]` to hold all the data.

Then, the code creates a `DNNClassifier` model using the following arguments:

- **feature\_columns=feature\_columns**. The set of feature columns defined above.
- **hidden\_units=[10, 20, 10]**. Three hidden layers (<http://stats.stackexchange.com/questions/181/how-to-choose-the-number-of-hidden-layers-and-nodes-in-a-feedforward-neural-netw>), containing 10, 20, and 10 neurons, respectively.
- **n\_classes=3**. Three target classes, representing the three Iris species.
- **model\_dir=/tmp/iris\_model**. The directory in which TensorFlow will save checkpoint data and TensorBoard summaries during model training.

## Describe the training input pipeline

The `tf.estimator` API uses input functions, which create the TensorFlow operations that generate data for the model. We can use `tf.estimator.inputs.numpy_input_fn` to produce the input pipeline:

```
# Define the training inputs
train_input_fn = tf.estimator.inputs.numpy_input_fn(
```

```
x={"x": np.array(training_set.data)},
y=np.array(training_set.target),
num_epochs=None,
shuffle=True)
```

## Fit the DNNClassifier to the Iris Training Data

Now that you've configured your DNN `classifier` model, you can fit it to the Iris training data using the `train` ([https://www.tensorflow.org/api\\_docs/python/tf/estimator/Estimator?hl=zh-cn#train](https://www.tensorflow.org/api_docs/python/tf/estimator/Estimator?hl=zh-cn#train)) method. Pass `train_input_fn` as the `input_fn`, and the number of steps to train (here, 2000):

```
# Train model.
classifier.train(input_fn=train_input_fn, steps=2000)
```

The state of the model is preserved in the `classifier`, which means you can train iteratively if you like. For example, the above is equivalent to the following:

```
classifier.train(input_fn=train_input_fn, steps=1000)
classifier.train(input_fn=train_input_fn, steps=1000)
```

However, if you're looking to track the model while it trains, you'll likely want to instead use a TensorFlow `SessionRunHook` ([https://www.tensorflow.org/api\\_docs/python/tf/train/SessionRunHook?hl=zh-cn](https://www.tensorflow.org/api_docs/python/tf/train/SessionRunHook?hl=zh-cn)) to perform logging operations.

## Evaluate Model Accuracy

You've trained your `DNNClassifier` model on the Iris training data; now, you can check its accuracy on the Iris test data using the `evaluate` ([https://www.tensorflow.org/api\\_docs/python/tf/estimator/Estimator?hl=zh-cn#evaluate](https://www.tensorflow.org/api_docs/python/tf/estimator/Estimator?hl=zh-cn#evaluate)) method. Like `train`, `evaluate` takes an input function that builds its input pipeline. `evaluate` returns a `dicts` with the evaluation results. The following code passes the Iris test data—`test_set.data` and `test_set.target`—to `evaluate` and prints the accuracy from the results:

```
# Define the test inputs
test_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={"x": np.array(test_set.data)},
    y=np.array(test_set.target),
    num_epochs=1,
    shuffle=False)

# Evaluate accuracy.
accuracy_score = classifier.evaluate(input_fn=test_input_fn) ["accuracy"]

print("\nTest Accuracy: {0:f}\n".format(accuracy_score))
```

**Note:** The `num_epochs=1` argument to `numpy_input_fn` is important here. `test_input_fn` will iterate over the data once, and then raise `OutOfRangeError`. This error signals the classifier to stop evaluating, so it will evaluate over the input once.

When you run the full script, it will print something close to:

```
Test Accuracy: 0.966667
```

Your accuracy result may vary a bit, but should be higher than 90%. Not bad for a relatively small data set!

## Classify New Samples

Use the estimator's `predict()` method to classify new samples. For example, say you have these two new flower samples:

Sepal Length	Sepal Width	Petal Length	Petal Width
6.4	3.2	4.5	1.5
5.8	3.1	5.0	1.7

You can predict their species using the `predict()` method. `predict` returns a generator of dicts, which can easily be converted to a list. The following code retrieves and prints the class predictions:

```
# Classify two new flower samples.
new_samples = np.array(
    [[6.4, 3.2, 4.5, 1.5],
     [5.8, 3.1, 5.0, 1.7]], dtype=np.float32)
predict_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={"x": new_samples},
    num_epochs=1,
    shuffle=False)

predictions = list(classifier.predict(input_fn=predict_input_fn))
predicted_classes = [p["classes"] for p in predictions]

print(
    "New Samples, Class Predictions:    {}\n"
    .format(predicted_classes))
```

Your results should look as follows:

```
New Samples, Class Predictions:    [1 2]
```

The model thus predicts that the first sample is *Iris versicolor*, and the second sample is *Iris virginica*.

## Additional Resources

- To learn more about using `tf.estimator` to create linear models, see [Large-scale Linear Models with TensorFlow](https://www.tensorflow.org/tutorials/linear?hl=zh-cn) (<https://www.tensorflow.org/tutorials/linear?hl=zh-cn>).
- To build your own Estimator using `tf.estimator` APIs, check out [Creating Estimators in tf.estimator](https://www.tensorflow.org/extend/estimators?hl=zh-cn) (<https://www.tensorflow.org/extend/estimators?hl=zh-cn>).
- To experiment with neural network modeling and visualization in the browser, check out [Deep Playground](http://playground.tensorflow.org/?hl=zh-cn) (<http://playground.tensorflow.org/?hl=zh-cn>).
- For more advanced tutorials on neural networks, see [Convolutional Neural Networks](https://www.tensorflow.org/tutorials/deep_cnn?hl=zh-cn) ([https://www.tensorflow.org/tutorials/deep\\_cnn?hl=zh-cn](https://www.tensorflow.org/tutorials/deep_cnn?hl=zh-cn)) and [Recurrent Neural Networks](https://www.tensorflow.org/tutorials/recurrent?hl=zh-cn) (<https://www.tensorflow.org/tutorials/recurrent?hl=zh-cn>).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](http://creativecommons.org/licenses/by/3.0/) (<http://creativecommons.org/licenses/by/3.0/>), and code samples are licensed under the [Apache 2.0 License](http://www.apache.org/licenses/LICENSE-2.0) (<http://www.apache.org/licenses/LICENSE-2.0>). For details, see our [Site Policies](https://developers.google.com/terms/site-policies?hl=zh-cn) (<https://developers.google.com/terms/site-policies?hl=zh-cn>). Java is a registered trademark of Oracle and/or its affiliates.

上次更新日期：十一月 2, 2017